

KOÇ UNIVERSITY  
Department of Computer Engineering  
Comp200 Structure and Interpretation of Computer Programs  
**The Evaluator**

- Submission: Perform a final save and copy the file to the following location  
F:/COURSES/UGRADS/COMP200/HOMEWORK/username/project5/project5.scm  
where *username* is your login name
- Code: The following code should be studied as part of this project:
  - `eval.scm`— code that defines the syntax, environments and control of the basic evaluator

As usual, you should begin working on the assignment once you receive it. It is to your advantage to get work done early, rather than waiting until the night before it is due. You should also read over and think through each part of the assignment (as well as any project code) before you sit down at the computer. It is generally much more efficient to test, debug, and run a program that you have thought about beforehand, rather than doing the planning "online." Diving into program development without a clear idea of what you plan to do generally guarantees that the assignments will take much longer than necessary.

The purpose of this project is to familiarize you with evaluators. We recommend that you first skim through the project to familiarize yourself with the format, before tackling problems.

Word to the wise: This project doesn't require a lot of actual programming. It does require understanding a body of code, however, and thus it will require careful preparation. You will be working with evaluators such as those described in chapter 4 of the text book, and variations on those evaluators. If you don't have a good understanding of how the evaluator is structured, it is very easy to become confused between the programs that the evaluator is interpreting, and the procedures that implement the evaluator itself. For this project, therefore, we suggest that you do some careful preparation. Once you've done this, your work on the computer should be fairly straightforward.

## 1 Understanding the evaluator

Load the code for this project. This file has three parts, and contains a version of the meta-circular evaluator similar to that described in lecture (there are a few minor differences) and in the textbook. The first part defines the syntax of the evaluator, the second part defines the actual evaluator, and the third part handles the environment structures used by the evaluator. Because this evaluator is built on top of the underlying Scheme evaluator, we have called the procedure that executes evaluation `m-eval` (with associated `m-apply`) to distinguish it from the normal `eval`.

You should look through these files to get a sense for how they implement a version of the evaluator discussed in lecture (especially the procedure `m-eval`).

You will be both adding code to the evaluator, and using the evaluator. Be careful, because it is easy to get confused. Here are some things to keep in mind:

When adding code to be used as part of `eval.scm`, you are writing in Scheme, and can use any and all of the procedures of Scheme. Changes you make to the evaluator are changes in defining the behavior you want your new evaluator to have.

After loading the evaluator (i.e., loading the file `eval.scm` and any additions or modifications you make), you start it by typing `(driver-loop)`. In order to help you avoid confusion, we've arranged that each driver loop will print prompts on input and output to identify which evaluator you are typing at. For example,

```
;;; M-Eval input:
```

```
(+ 3 4)
```

```
;;; M-Eval value:
```

```
7
```

shows an interaction with the `m-eval` evaluator. To evaluate an expression, you type the expression to showed up box and hit enter.

The evaluator with which you are working does not include an error system. If you hit an error you will bounce back into ordinary Scheme. You can restart the driver-loop by running the procedure `(driver-loop)`. Note that the driver loop does not re-initialize the environment, so any definitions you have made should still be available if you have to re-run `driver-loop`. In case you do want a clean environment, you should evaluate `(refresh-global-environment)` while in normal Scheme.

To quit out of the new evaluator, simply evaluate the expression `**quit**`. This will return you to the underlying Scheme evaluator and environment.

## 1.1 Computer Exercise 1: Exploring Meval

Load the code files into your Scheme environment. To begin using the interpreter defined by this file, evaluate `(driver-loop)`. Notice how it now gives you a new prompt, in order to alert you to the fact that you are now "talking" to this new interpreter. Try evaluating some simple expressions in this interpreter.

You will probably very quickly notice that some of the primitive procedures about which Scheme normally knows are missing in `m-eval`. These include some simple arithmetic procedures (such as `*`) and procedures such as `cadr`, `cddr`, `newline`, `length`. Extend your evaluator by adding these new primitive procedures (and any others that you think might be useful). Check through the code to figure out where this is done. In order to make these changes visible in the evaluator, you'll need to rebuild the global environment:

```
(refresh-global-environment)
```

or you will need to re-evaluate your file (click the run button) and start fresh (which is probably the less confusing option). Show your changes to the evaluator, and turn in a demonstration of your extended evaluator working correctly.

## 2 Adding new special forms to Scheme

### 2.1 Computer Exercise 2: Changing style

In our standard evaluator, if we `define` a variable with a value for a second time, we lose the previous value of the variable. Similarly, if we `set!` a variable, we lose the previous value. We

would like to change this behavior, so that we keep track of all previous bindings of a variable. This would allow us to introduce a new **special form**, `reset!`. Evaluating `(reset! var)` in some environment should have the following behavior:

- if the variable `var` had been previously defined, its value is reset to the most recent prior value,
- if the variable `var` had only been defined once, then this evaluation results in an error (since there is no error mechanism, just display the error message: "ERROR: No value to reset!"),
- `reset!` can be used multiple times, that is, evaluating `(reset! var) (reset! var)` will "undo" the two previous bindings for `var` (assuming that it had been set more than twice).

To implement this change in our evaluator, we need to do several things:

- Change bindings of variables to associate a list of values with a variable, rather than a single value (note that you can do this by simply changing `set-binding-value!`).
- Add a new special form called `reset!`. Evaluating this special form should change the binding of the variable to include the new value in the list of bindings that have been associated with the variable. (You may find `lookup-variable-value` to be a useful template for this change.) Be sure to think about where a special form should go in `m-eval` as well as creating a syntactic abstraction to handle `reset!` expressions.

Implement this change and demonstrate your new evaluator showing this new behavior.

## 2.2 Computer Exercise 3: Adding a new special form

We have seen that our evaluator treats any compound expression as an application of a procedure, unless that expression is a special form that has been explicitly defined to obey a different set of rules of evaluation (e.g., `define`, `lambda`, `if`). We are going to add a special form to our evaluator in the next exercise.

A common construct in other languages is a loop. Our version of a loop has the following syntax:

```
(loop <until> <return>
  <exp1>
  <exp2>
  ...
  <expN>)
```

The behavior is as follows: Each of the `exp1 ... expN` expressions is evaluated in order. Then the `until` expression is evaluated. If the `until` expression evaluates to `#t` then we stop executing the loop and return the value obtained by evaluating the `return` expression. Otherwise we evaluate the expressions `exp1` through `expN` again, and repeat the whole process. In short, a loop repeatedly evaluates `exp1`, `exp2`, ..., `expN` until the expression `until` evaluates to `#f`, whereupon the entire loop expression evaluates to the value of the `return` expression.

Here's an example:

```

(let ((x '()))
  (loop (> (length x) 3) x
    (set! x (cons '* x))
    (display x)))

(*) (* *) (* * *) (* * * *)

; M-eval value: (* * * *)

```

Your task is to add this special form to `m-eval`, showing your changes to the code, and demonstrating that it works by providing some test cases.

To do this, you should include the following:

- Create a data abstraction for handling `loops`, that is, selectors for getting out the parts of the loop that will be needed in the evaluation.
- Add a dispatch clause to the correct part of `m-eval`.
- Write the procedure(s) that handle the actual evaluation of the `loop`.

Be sure to turn in a listing of your additions and changes, as well as examples of your code working on test cases.

### 3 Transformers: More than meets the eye

In some cases, it is easy to think of implementing a special form in terms of more primitive expressions that are already covered by the evaluator. As an example, to evaluate `cond` expressions, `m-eval` uses the following code in the dispatcher:

```
((cond? exp) (m-eval (cond->if exp) env))
```

together with the following method for converting a `cond` statement to an `if` statement:

```

(define (cond->if expr)
  (let ((clauses (cond-clauses expr)))
    (if (null? clauses)
        #f
        (if (eq? (car (first-cond-clause clauses)) 'else)
            (make-begin (cdr (first-cond-clause clauses)))
            (make-if (car (first-cond-clause clauses))
                     (make-begin (cdr (first-cond-clause clauses)))
                     (make-cond (rest-cond-clauses clauses)))))))

```

What would the following expression evaluate to?:

```

(cond->if '(cond ((= x 0) 5)
              (> x 0) 10)
          (else 15)))

```

### 3.1 Computer Exercise 4: Transforming boolean combinations

In scheme, `and` and `or` are special forms because all of the arguments to these forms are not necessarily evaluated. The following example exhibits this behaviour:

```
(define (safe-list-ref lst n)
  (if (and (integer? n) (list? lst) (>= n 0) (< n (length lst)))
      (list-ref lst n)
      'invalid-list-reference))
```

So, for example, if `(list? lst)` is false, then the expression `(length lst)` is not evaluated, which is good. This is called 'short-circuiting:' as soon as an argument to `and` evaluates to false we return true, without evaluating the remaining arguments; `or` behaves almost the same way. You already implemented this behavior in the problem set, but there's one more feature we want you to add.

The result of an `or` special form is not necessarily `#t` or `#f`, but rather the value of the first non-`#f` argument. (For example, `(or #f 3 4) => 3`.) Likewise, if an `and` special form has no arguments which evaluate to `#f`, then its return value is the value of its last argument. (For example, `(and 2 3 4) => 4`.)

As before, add a pair of clauses to `m-eval` that transform `or` and `and` expressions into expressions that `m-eval` already knows how to evaluate. Remember that the individual expressions in the body of `or` are supposed to be evaluated at most once:

```
(or (begin (display "once ") #f)
    (begin (display "and ") #f)
    (begin (display "only ") 'done)
    (begin (display "adbm1") #t))
```

```
once and only
;Value: done
```

```
(and (begin (display "a ") (> 6 5))
      (begin (display "b ") (< 6 5))
      (begin (display "c ") (+ 6 5)))
a b
;Value: #f
```

(One way to desugar `or` involves using `let` and `if`. In this case there can be "name capture" errors if the `let` variable already appears in the `or` expression being desugared. For this problem, it is okay to assume that any variables you use in desugaring do not already appear. Note that there is a better way to desugar `and` and `or` that avoids this problem.)

### 3.2 Computer Exercise 5: Implementing `loop2` as syntactic sugar

Show how to implement `loop2` using a similar strategy to that for `cond`, i.e., convert a `loop2` expression into an associated `if` expression. You should have the following code in the dispatcher:

```
((loop2? exp) (m-eval (loop->if exp) env))
```

together with a definition for `loop->if` (hint: `loop->if` may return an expression which includes a `loop2` expression as one of its sub-expressions).

## 4 Implementing an Evaluator with Memoization

A few times in the course we've seen the idea of *memoization*. This involves using some kind of data structure to store previously computed values of a procedure application.

In this part of the project we're going to extend the evaluator to allow automatic implementation of memoization. In particular, we are going to extend the syntax of lambda expressions so that they allow us to specify if we want to use memoization with the compound procedure created by the lambda expression.

The new syntax will be as follows:

```
(lambda <memo?> <parameters> <body>)
```

The `<parameters>` and `<body>` are the same as in regular lambda expressions. `<memo?>` is either `'memo` or `'no-memo`, and specifies whether or not memoization is desired with this lambda expression.

As an example, we might do the following:

```
(define myplus (lambda 'memo (x y) (+ x y)))

(myplus 3 4)    ;-> return value is 7
(myplus 3 4)    ;-> return value is 7
```

In the first call to `myplus` with the arguments 3 and 4, the procedure operates in its usual way. In the second call, however, the evaluator realizes that the arguments 3 and 4 have been used before with this procedure, and the value 7 is retrieved from a data structure that stores the result of previous computations with `myplus`.

On the other hand, if we do the following:

```
(define myplus2 (lambda 'no-memo (x y) (+ x y)))

(myplus2 3 4)   ;-> return value is 7
(myplus2 3 4)   ;-> return value is 7
```

then memoization is not used (because `'no-memo` was specified in the original lambda expression), and in both cases the value 7 has to be computed from scratch.

Notice that this use of memoization can be very useful in some cases, take for example the following definition of `fibonacci`:

```
(define fib (lambda 'memo (n)
  (cond ((= n 1) 1)
        ((= n 2) 1)
        (else (+ (fib (- n 1)) (fib (- n 2)))))))
```

This will immediately give an implementation of `fibonacci` with memoization; this procedure should run in  $\Theta(n)$  time, rather than the exponential time required for a version without memoization.

As one additional point, note that we'll have to alter the syntax for `define` statements as well. In regular scheme, we can define procedures in the following way:

```
(define (myplus x y) (+ x y))
```

In the new version with memoization, in this case we could use either

```
(define (myplus 'memo x y) (+ x y))
```

or

```
(define (myplus 'no-memo x y) (+ x y))
```

to define versions of `myplus` with or without memoization.

To implement memoization, a crucial idea will be the following: *a lambda expression evaluates to a triple bubble in the new evaluator*. The first two bubbles are the same as in regular scheme: the first bubble points to the parameters and body of the lambda expression, the second bubble points to the environment in which the lambda expression was evaluated. The third bubble points to a data structure if the lambda expression had the `'memo` option, or the symbol `none` if it had the `'no-memo` option. Initially the data structure is empty.

You will need to implement a data structure to store information about previous calls to a procedure. One easy way to do this is to use a list, where each entry in the list is itself a list consisting of the values supplied to the procedure and the associated value. For example, if `fib` had been called with the arguments 4, 3, 2 and 1, we **might** choose to represent that information as:

```
((5) 5) ((4) 3) ((3) 2) ((2) 1) ((1) 1))
```

Note that the strange form of this list is because a procedure might take more than one argument, so we are capturing in each element a list of the list of arguments and the value.

When using `m-apply`, for compound procedures a triple bubble will be applied to a list of arguments. An application of a triple bubble involves the following steps:

- If the third bubble points to `none`, then apply the compound procedure in the usual way (no memoization is needed).
- Otherwise, the third bubble points to a memoization data structure. In this case, first check to see if the `arguments` are stored in the data structure with some value. If so, return that value as the value of the application. Otherwise, evaluate the compound procedure in the usual way, *but make sure to store an entry in the data structure that associates the arguments with the value that is computed, before returning that value.*

## 4.1 Computer Exercise 6: Adding memoization to the evaluator

To implement this style of memoization in the mc-evaluator, you'll need to carry out the following steps:

- Change and add to the code defining the syntax of lambda expressions, to take into account the modification that allows the additional `'memo` or `'no-memo` syntax. (Note that you should change `make-lambda` to take an additional argument, `memo?`, and that the call to `make-lambda` within `let->application` should use the `'no-memo` option.)
- Change the code for the `define` syntax to take into account the additional `'memo` or `'no-memo` option.

- Create a data structure for storing memoization information about application of a procedure – what values of arguments has it been called on, and what was the resulting value.
- Change the code for `make-procedure` so that it takes an additional argument that is either `'memo` or `'no-memo`. Add a function `procedure-m-list` which returns a memoization data structure from the third bubble if the `'memo` option was used, or the symbol `none` otherwise. Change the call to `make-procedure` in the dispatcher.
- Change `m-apply` to implement memoization. See the description above for how this should procede.

Be sure to turn in a listing of your additions and changes, as well as examples of your code working on test cases.

Finally, note that code with and without memoization may give different values, in particular when mutation or side-effects are used. Give three examples of code that gives different behavior (i.e., different return values, or different output to the terminal) when the `'memo` and `'no-memo` options are used.

## 5 Submission

For each problem, include your code (with identification of the exercise number being solved), as well as comments and explanations of your code, and demonstrate your code's functionality against a set of test cases. Once you have completed this project, your file should be submitted electronically according to this handout.

We encourage you to work with others on problem sets as long as you acknowledge it (nothing written gets exchanged!).

If you cooperated with other students, TA's, or tutors, please indicate your consultants' names and how they collaborated. Be sure that your actions are consistent with the posted course policy on collaboration.